

# Case Study 1

Jonathan Sneh, Ishani Tarafdar, Georges Durand, Raul Higareda

2023-03-28

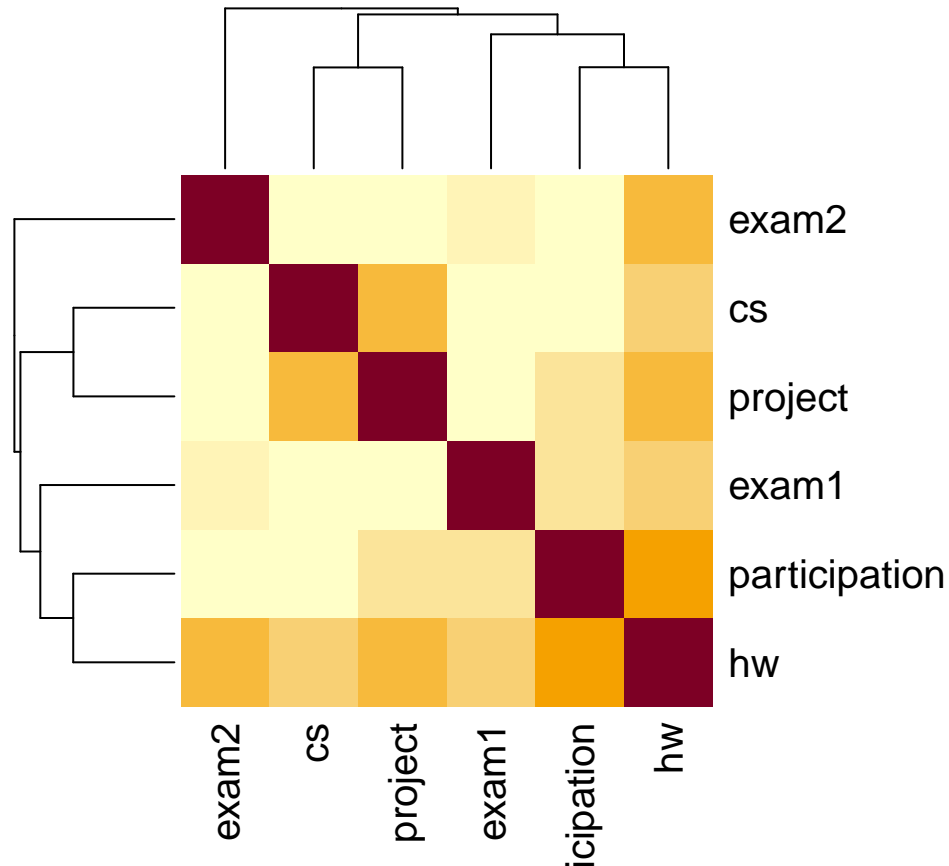
## Data Explorations and Summary Statistics

```
grades <- read.csv("grades.csv", header=TRUE)
grades <- grades[1:6]
dim(grades)
```

```
## [1] 275 6
```

## Model Selection

```
cor_matrix <- cor(grades)
heatmap(cor_matrix, symm = TRUE)
```



Looking at the correlation matrix, none of the variables are strongly correlated with another variable, so we do not need to drop any variables based on what we see from the correlation matrix.


We want to make a model with 95% confidence (i.e.  $\alpha = 0.05$ )

```
grades.mlr <- lm(exam2 ~ hw + cs + participation + exam1 + project ,data=grades)
summary(grades.mlr)
```

```
##
## Call:
## lm(formula = exam2 ~ hw + cs + participation + exam1 + project,
##     data = grades)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.116  -9.270   3.431  10.486  39.096
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  25.51074     6.52460   3.910 0.000117 ***
## hw           0.50233     0.07506   6.692 1.27e-10 ***
## cs           0.02096     0.07030   0.298 0.765805
## participation -0.09622     0.06259  -1.537 0.125403
## exam1        0.22437     0.06954   3.226 0.001410 **
## project      0.00885     0.04804   0.184 0.853973
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
##  
## Residual standard error: 14.59 on 269 degrees of freedom  
## Multiple R-squared:  0.3249, Adjusted R-squared:  0.3124  
## F-statistic: 25.9 on 5 and 269 DF,  p-value: < 2.2e-16
```

```
hist(grades.mlr$residuals)
```





```
predicted_values <- predict(grades.mlr)  
plot(x = predicted_values, y = residuals(grades.mlr),  
     xlab = "Predicted values", ylab = "Residuals",  
     main = "Residual plot for heteroscedasticity check")
```


### Residual plot for heteroscedasticity check




```
plot(grades.mlr)
```









Based on the full model summary above, we may want to look into dropping project, cs, and participation from the dataset since the t-values in the summary output for project, cs, and participation are all small, meaning that they're likely up to chance.

However, the individual t-tests do not tell us enough information to drop multiple predictors from our model at a time.

So, we can start by dropping an individual predictor from our model. We will check **project**.

Our null and alternative hypothesis are as follows.

$$\begin{cases} H_0, & \beta_{\text{project}} = 0 \\ H_A, & \beta_{\text{project}} \neq 0 \end{cases}$$

By conducting an individual t-test (which can be found in our summary output), we can see that the t-test statistics for project is 0.397. **## TODO PLEASE CHECK IF WE NEED TO BE THOROUGH AND MENTION WHICH DISTRIBUTION THESE ARE FROM** This comes from a  $t$ . Thus, we can see that the p-value for project is 0.692.  $p = 0.692 > 0.05 = \alpha$ . Thus, we fail to reject the null hypothesis, meaning that it is likely that  $\beta_{\text{project}} = 0$ . In other words, we can drop project from our model.

This leaves us with the following reduced model:

```
grades.reducedmlr1 = lm(exam2~exam1 + hw + cs + participation, data=grades)
summary(grades.reducedmlr1)
```

```
##
## Call:
```



```
## lm(formula = exam2 ~ exam1 + hw + cs + participation, data = grades)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.129  -9.371   3.396  10.476  39.564
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  25.36389     6.46413   3.924 0.000111 ***
## exam1         0.22444     0.06942   3.233 0.001377 **
## hw           0.50605     0.07216   7.013 1.88e-11 ***
## cs           0.02568     0.06535   0.393 0.694589
## participation -0.09430     0.06161  -1.531 0.127027
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 14.56 on 270 degrees of freedom
## Multiple R-squared:  0.3249, Adjusted R-squared:  0.3148
## F-statistic: 32.48 on 4 and 270 DF,  p-value: < 2.2e-16
```

```
anova(grades.mlr, grades.reducedmlr1)
```

```
## Analysis of Variance Table
##
## Model 1: exam2 ~ hw + cs + participation + exam1 + project
## Model 2: exam2 ~ exam1 + hw + cs + participation
##   Res.Df  RSS Df Sum of Sq    F Pr(>F)
## 1      269 57240
## 2      270 57247 -1    -7.2221 0.0339  0.854
```

Can we reference the overall summary and p-values to give us guidance on where to look next?

From the summary, We can see that the p-values for cs and participation have changed. They are still high, so we can conduct a different test.

$$\begin{cases} H_0, & \beta_{participation} = \beta_{cs} = 0 \\ H_A, & \text{Either } \beta_{participation} \text{ or } \beta_{cs} \text{ is not equal to zero} \end{cases}$$

```
library(ellipse)
```

```
## Warning: package 'ellipse' was built under R version 4.1.3
```

```
##
## Attaching package: 'ellipse'
```

```
## The following object is masked from 'package:graphics':
##
##      pairs
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.1.3
```


We can draw the confidence region (as an ellipse) for both participation and for cs. If the point (0,0) falls inside of our confidence region, then it is likely that both the coefficients  $\beta_{cs}$  and  $\beta_{participation}$  are zero—as according to the null hypothesis.

```
intervals <- confint(grades.reducedmlr1)
cr_ellipse <- ellipse(grades.reducedmlr1, c(4,5), level=0.95)

par_interval <- confint(grades.reducedmlr1, level = 0.95, 'participation')
cs_interval <- confint(grades.reducedmlr1, level = 0.95, 'cs')

cr_df <- as.data.frame(cr_ellipse)
cr_plot <-
  ggplot(data=cr_df, aes(x=participation, y=cs)) +
    ggtitle("Confidence Region -- Joint Estimation for participation and cs") +
    geom_path(aes(x=participation,y=cs), colour='mediumorchid') +
    geom_point(x=coef(grades.reducedmlr1)[2], y=coef(grades.reducedmlr1)[3],
              shape=3, size=3, colour='mediumorchid') +
    geom_hline(yintercept = cs_interval[1], lty=2) +
    geom_hline(yintercept = cs_interval[2], lty=2) +
    geom_vline(xintercept = par_interval[1], lty=2) +
    geom_vline(xintercept = par_interval[2], lty=2)+
    geom_point(x=0, y=0, shape=1, size=3, colour='green')

plot(cr_plot)
```



As we can see, the origin—which is the green dot—falls inside the confidence region. Thus, it is likely enough that both  $\beta_{cs}$  and  $\beta_{participation}$  are zero. Therefore, we can drop them both from our model.

Since the p-value in the full model is the same as the reduced model, this confirms that the dropped variables were not significant in our model. The sample size is quite small, therefore dropping variables may lead to insignificant changes in p value.

So, our final model (before diagnostics) is the following:

```
grades.reducedmlr = lm(exam2 ~ exam1 + hw, data=grades)
summary(grades.reducedmlr)
```

```
##
## Call:
## lm(formula = exam2 ~ exam1 + hw, data = grades)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -51.150  -9.195   3.301  10.349  40.987
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.64208    5.61643   4.387 1.64e-05 ***
## exam1       0.20858    0.06780   3.077 0.00231 **
## hw          0.46396    0.05768   8.044 2.69e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 14.57 on 272 degrees of freedom
## Multiple R-squared:  0.3187, Adjusted R-squared:  0.3137
## F-statistic: 63.62 on 2 and 272 DF,  p-value: < 2.2e-16
```


## Unusual Observations and Model Assumptions

Now we can analyze the final model for unusual observations and check for deviations from the model assumptions.

### Constant variances

First, we can check the model assumption for constant variances by checking the residual vs. fitted plot.

```
plot(grades.reducedmlr, which=1)
```




From the residuals vs. fitted plots, we can see the the assumptions for constant variance are not met because the residuals are not evenly distributed around the 0 line, and seem to decrease in magnitude as the residuals increase.

### Normality


Next, we can ckck for normaltiy of the residuals by creating a QQ plot.

```
plot(grades.reducedmlr, which=2)
```



```
predicted_values <- predict(grades.reducedmlr)
plot(x = predicted_values, y = residuals(grades.reducedmlr),
     xlab = "Predicted values", ylab = "Residuals",
     main = "Residual plot for heteroscedasticity check")
```

## Residual plot for heteroscedasticity check




From the QQ plot, we can see that we seem to have departures from the normality assumption as points along the edges of the plot don't follow the straight line. We can attempt to remedy this and reduce the non-normality of the errors by performing a Box-Cox transformation.

```
library(MASS)
```

```
## Warning: package 'MASS' was built under R version 4.1.3
```

```
grades.transformation = boxcox(grades.reducedmlr, lambda=seq(-2, 2, length=400))
```



Looking at the Box-Cox plot, the optimal  $\lambda$  is somewhere between 1.5 and 2. We can try both to see which provides better results.

```
grades2 = grades
grades2$grades.new = ('^'(grades2$exam2,2)-1) / 2
grades.mlr.tr = lm(grades.new ~ exam1 + hw, data=grades2[])
summary(grades.mlr.tr)
```

```
##
## Call:
## lm(formula = grades.new ~ exam1 + hw, data = grades2[])
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-3074.2	-794.3	175.5	869.6	2702.0

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-526.872	415.703	-1.267	0.206
exam1	15.659	5.018	3.120	0.002 **
hw	31.640	4.269	7.412	1.58e-12 ***


```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1079 on 272 degrees of freedom
## Multiple R-squared:  0.2932, Adjusted R-squared:  0.2881
## F-statistic: 56.43 on 2 and 272 DF,  p-value: < 2.2e-16
```


```
plot(grades.mlr.tr)
```











### Serial Dependence

It is not possible to check serial dependence for this model because there is no order or time value associated with the data points.

### Unusual Observations

```
grades.leverages = lm.influence(grades.reducedmlr)$hat
head(grades.leverages)
```

### High Leverage points

```
##          1          2          3          4          5          6
## 0.007424924 0.007274227 0.006206353 0.006622878 0.005729212 0.005620944
```

```
n = dim(grades)[1]
p = length(variable.names(grades.reducedmlr))

grades.leverages.high = grades.leverages[grades.leverages > 2*p/n]
grades.leverages.high
```

```
##          33          34          39          45          47          50          51
## 0.02272018 0.02206615 0.02583001 0.03926475 0.09845203 0.03631088 0.05592355
##          138          153          159          160          162          164          165
## 0.04149202 0.02290199 0.03202090 0.03833443 0.02218738 0.02563467 0.02268266
##          166          167          168          169          250          253          254
## 0.08057361 0.03358242 0.08762130 0.07687729 0.02502428 0.07983386 0.13633458
##          273          274          275
## 0.02469497 0.02964430 0.08837007
```

There are 23 high leverage points within our reduced model versus 50 in our full model. These points have the potential of increasing variability, so having less in our reduced is ideal. This also shows how the dropped variables contributed to high leverage points in the full model.

```
grades.leverages = lm.influence(grades.reducedmlr)$hat
head(grades.leverages)
```

## Outliers

```
##          1          2          3          4          5          6
## 0.007424924 0.007274227 0.006206353 0.006622878 0.005729212 0.005620944
```

```
n = dim(grades)[1]
p = length(variable.names(grades.reducedmlr))

# Computing Studentized Residuals
grades.resid = rstudent(grades.reducedmlr);

# Critical value WITH Bonferroni correction
bonferroni_cv = qt(.05/(2*n), n-p-1)

grades.resid.sorted = sort(abs(grades.resid), decreasing=TRUE)[1:10]
print(grades.resid.sorted)
```


```
##          151          272          250          47          247          227          241          45
## 3.601063 3.404720 3.105902 3.005424 2.962759 2.462233 2.371182 2.263594
##          253          226
## 2.248131 2.229855
```

```
grades.outliers = grades.resid.sorted[abs(grades.resid.sorted) > abs(bonferroni_cv)]
print(grades.outliers)
```

```
## named numeric(0)
```


```
hist(grades.reducedmlr$residuals)
```

**Histogram of grades.reducedmlr\$residuals**



```
hist(grades.mlr$residuals)
```

## Histogram of grades.mlr\$residuals




```
grades.cooks = cooks.distance(grades.reducedmlr)
sort(grades.cooks, decreasing = TRUE)[1:10]
```

### Influential Observations

```
##          47          253          250          45          275          166          51
## 0.31936425 0.14401823 0.07998918 0.06876062 0.05806449 0.05763882 0.04667620
##          169          138          151
## 0.04581583 0.03546309 0.03443911
```


```
plot(grades.cooks)
```



```
library(faraway)
```

```
## Warning: package 'faraway' was built under R version 4.1.3
```

```
halfnorm(grades.cooks, 6, labs=as.character(1:length(grades.cooks)), ylab="Cook's distances")
```




```
grades.cooks = cooks.distance(grades.mlr)
sort(grades.cooks, decreasing = TRUE)[1:10]
```


```
##          47          253          51          45          166          250          167
## 0.27837732 0.15771394 0.08485803 0.05961187 0.04761373 0.04711742 0.04252647
##          153          275          169
## 0.03816061 0.03492719 0.03492200
```

```
plot(grades.cooks)
```





```
library(faraway)
halfnorm(grades.cooks, 6, labs=as.character(1:length(grades.cooks)), ylab="Cook's distances")
```



The number of influential points between the reduced and full model remained the same. It is not bad because the number of points did not increase.